

## Q. 4.2.1

There is a lot of data redundancy in the given table, as it hasn't been normalised. This is not a very efficient way to store data. Without any normalisation of a relational database, we would create giant tables that have a lot of repetitive data and take up a large amount of memory. By increasing the complexity of the data model, and partitioning the data into separate tables, we increase the efficiency of the database.

Mary, Joe and Tom are entered twice each into the StudentName columns. To improve the efficiency of the data storage, I would partition the StudentName column, along with the dob column to a separate table, to have a unique record per student. This table would have a 1:1 relationship. See figure 1 below.

/* Student table */		
studentID (primary key)	studentName	dob
1	Sean	2000-01-03
2	Bill	1990-04-23
3	Tom	1973-12-10
4	Mary	1991-04-12
5	Joe	1982-06-29

Figure 1

There is also a lot of repetition/redundancy of data in the moduleID and moduleName columns of the original table. I would also separate out these columns into another 1:1 table with unique/distinct values (see example table in figure 2), and then reference them in a 3rd (1:many) table that combines the studentIDs with the moduleIDs (see example table in figure 3). The primary keys of both the Module table and the Student table are referenced in the StudentModules table as foreign keys.

/* Module table */	
moduleID (primary key)	moduleName
100	Applied Databases
101	Java Programming
102	Mobile Apps
103	Computer Architecture

Figure 2

/* StudentModules table */		
studentModuleID (primary key)	moduleID (foreign key)	studentID (foreign key)
1	100	1
2	100	2
3	101	3
4	104	3
5	101	4
6	102	4
7	101	5
8	104	5

Figure 3

A join query can be done between the primary keys of the Student table, the Module table and their foreign keys in the StudentModule table to find which students are taking which courses, thus reducing data redundancy in the data model. See an example design of the new data model ERD below in figure 4.

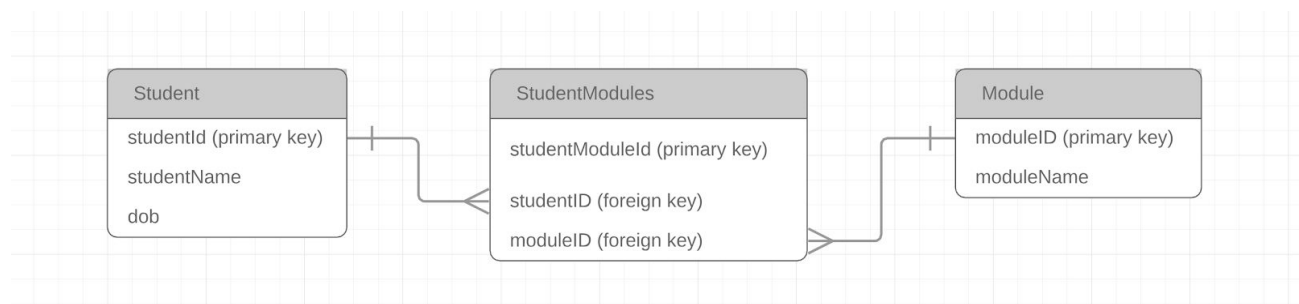


Figure 4

The above data model design (figure 4) follows the 3<sup>rd</sup> normal form, which in most practical applications, is the most appropriate form of normalisation, and adheres to the following design rules: (1)(2)

- Each table cell should contain a single value (1<sup>st</sup> normal form).
- Each record in each table is unique (1<sup>st</sup> normal form).
- Single column primary keys (2<sup>nd</sup> normal form) - Users will only be able to insert values into a foreign key field that exist as a primary key in the parent table ie: we cannot insert a new studentID in the StudentModules table if it doesn't exist in the Student table. This helps to maintain referential integrity.
- There are no transitive functional dependencies (3<sup>rd</sup> normal form) – if we update a non-key column, it won't cause data in the other columns to change.

#### References:

- 1 Guru99, 'What is Normalization? 1NF, 2NF, 3NF & BCNF with Examples', accessed on 10<sup>th</sup> May 2019, <<https://www.guru99.com/database-normalization.html>>
- 2 Codd, E. F., 'Relational Database: A Practical Foundation for Productivity', Communications

